

Distributed Computation with Python and Celery

Hi

Lars.Butler@gmail.com

www.openquake.org

Overview

Part 1:

- What is Celery?
- How is it useful?
- How does it work?

Part 2:

- What is a 'task'?
- What are the supported message brokers?
- How are task results stored?
- Configuring Celery

Demo:

- Architecture examples
- Code examples

Questions

Part 1

Conceptual stuff

What is Celery?

- A delicious vegetable
- An asynchronous task queue based on distributed messaging

Homepage: <http://celeryproject.org>

Source code: <http://pypi.python.org/pypi/celery#downloads>

Revision control: <https://github.com/ask/celery>

License: BSD-style (<https://github.com/ask/celery/blob/master/LICENSE>)

Documentation: <http://docs.celeryproject.org/>


IRC: #celery on irc.freenode.net

Version 1.0: February 10th 2010

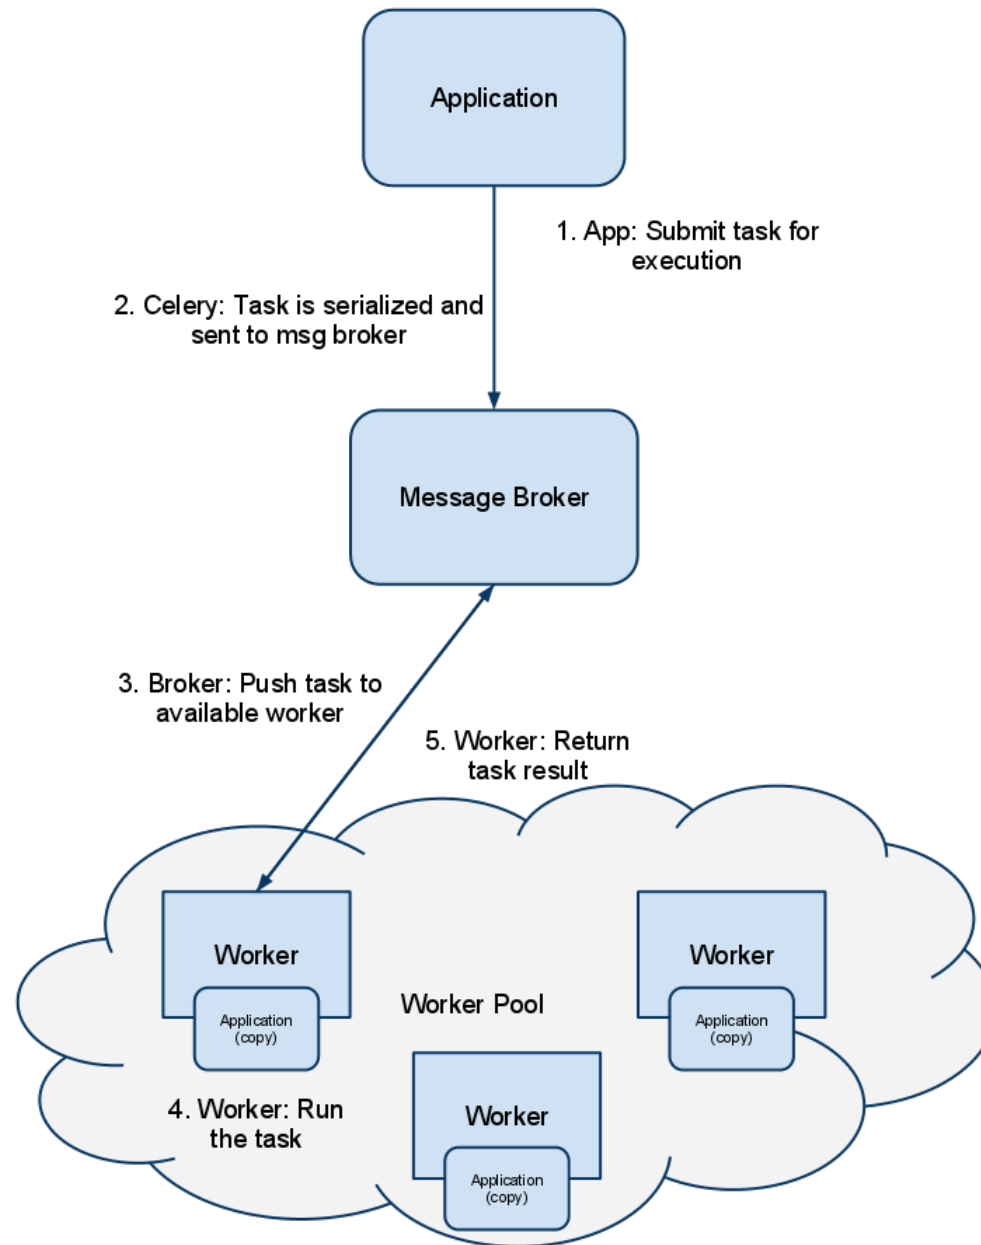
Version 2.2: February 1st 2011 (latest version)

How is Celery useful?

- Break computations down into small chunks ('tasks')
- Process tasks using remote workers
- Encapsulates horizontal scaling concerns



How does it work?



Part 2

Details

What is a 'task'?

- A Celery task is a set of computations
 - Can be executed asynchronously
 - Can be purely procedural or return a result
- Tasks are functions or methods with the '@task' decorator from the celery.decorators package
- @task must be the 'outermost' decorator, like so:

```
@task
```

```
@foo
```

```
@bar
```

```
def some_task(x, y):
```

```
    # task code here
```

Example tasks

- Complex/time-consuming mathematical computations
- Image processing/rendering
- Searching/sorting large data sets
- Pathfinding algorithms on large data sets
- Servicing a client request in a high-availability system
- Virtually anything that requires a non-trivial amount of time or processing power

What are the supported message brokers?

Preferred:

- RabbitMQ

Limited support:

- Redis
- Beanstalk
- MongoDB
- CouchDB
- RDBMSs (using SQLAlchemy or Django ORM)
- Others

How are task results stored?

- Result backends supported:
 - database (default)
 - MongoDB
 - Redis
 - AMQP
 - Others
- If task result data is of a non-trivial size, it may be useful to store the data produced by the task in a database or key value store.

Configuring Celery: AMQP

Example celeryconfig.py:

```
import sys
```

```
sys.path.append('.')
```

```
BROKER_HOST = "localhost"
```

```
BROKER_PORT = 5672
```

```
BROKER_USER = "celeryuser"
```

```
BROKER_PASSWORD = "celery"
```

```
BROKER_VHOST = "celeryvhost"
```

```
CELERY_RESULT_BACKEND = "amqp"
```

```
CELERY_IMPORTS = ("tasks",)
```

Configuring Celery: SQLite

```
import sys

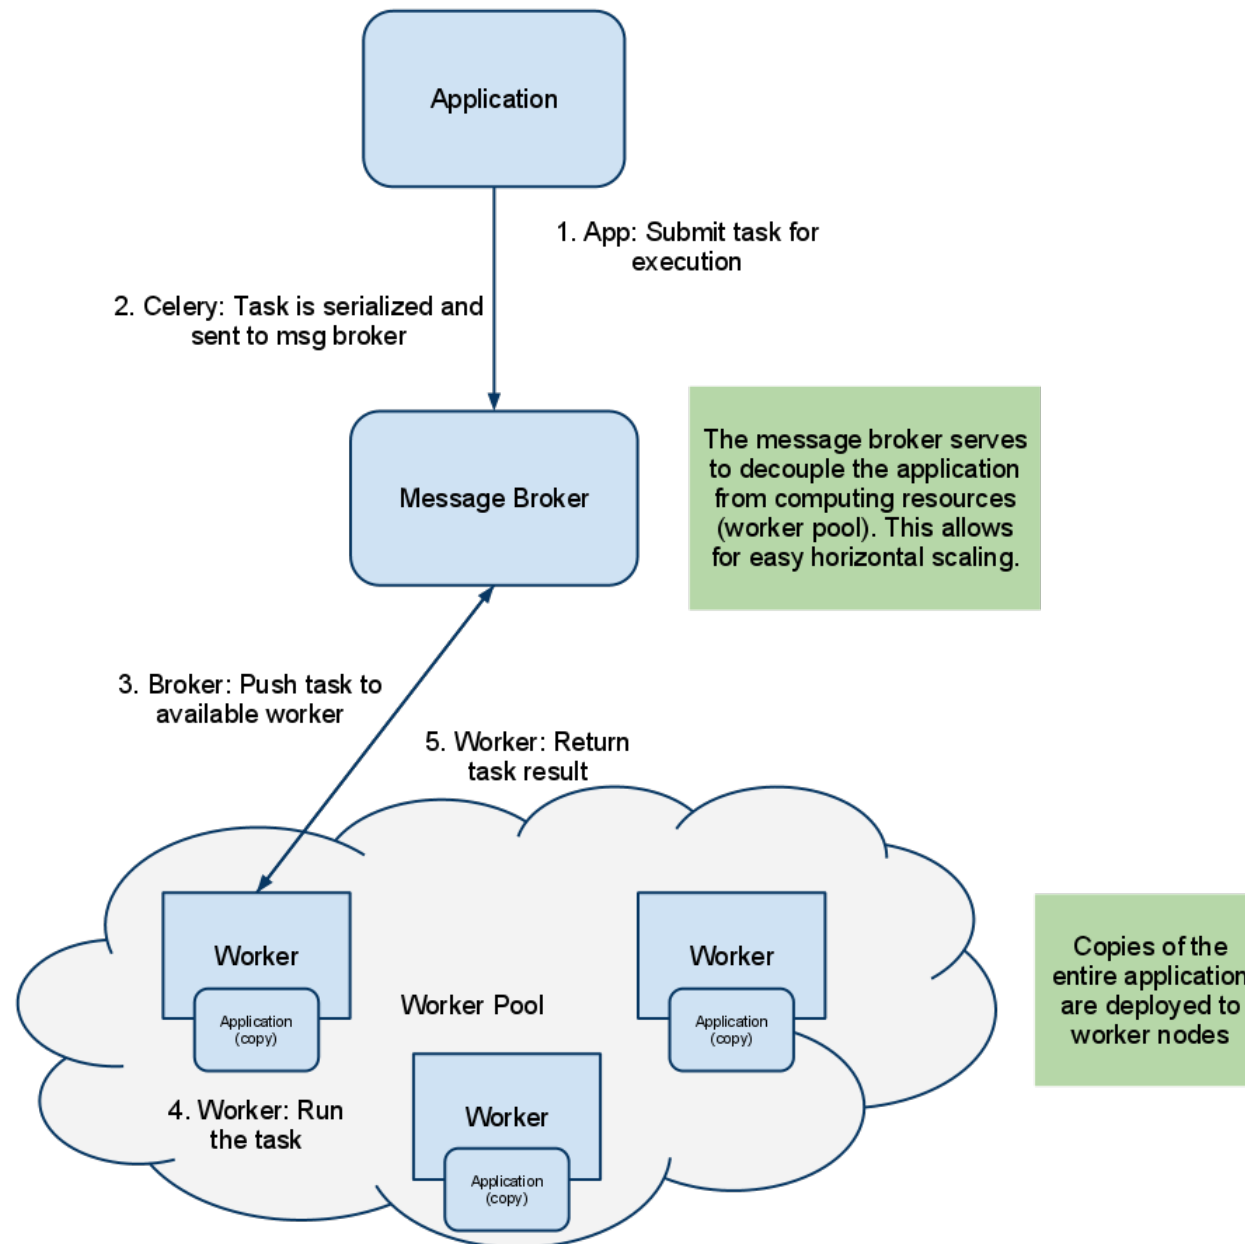
sys.path.append('.')

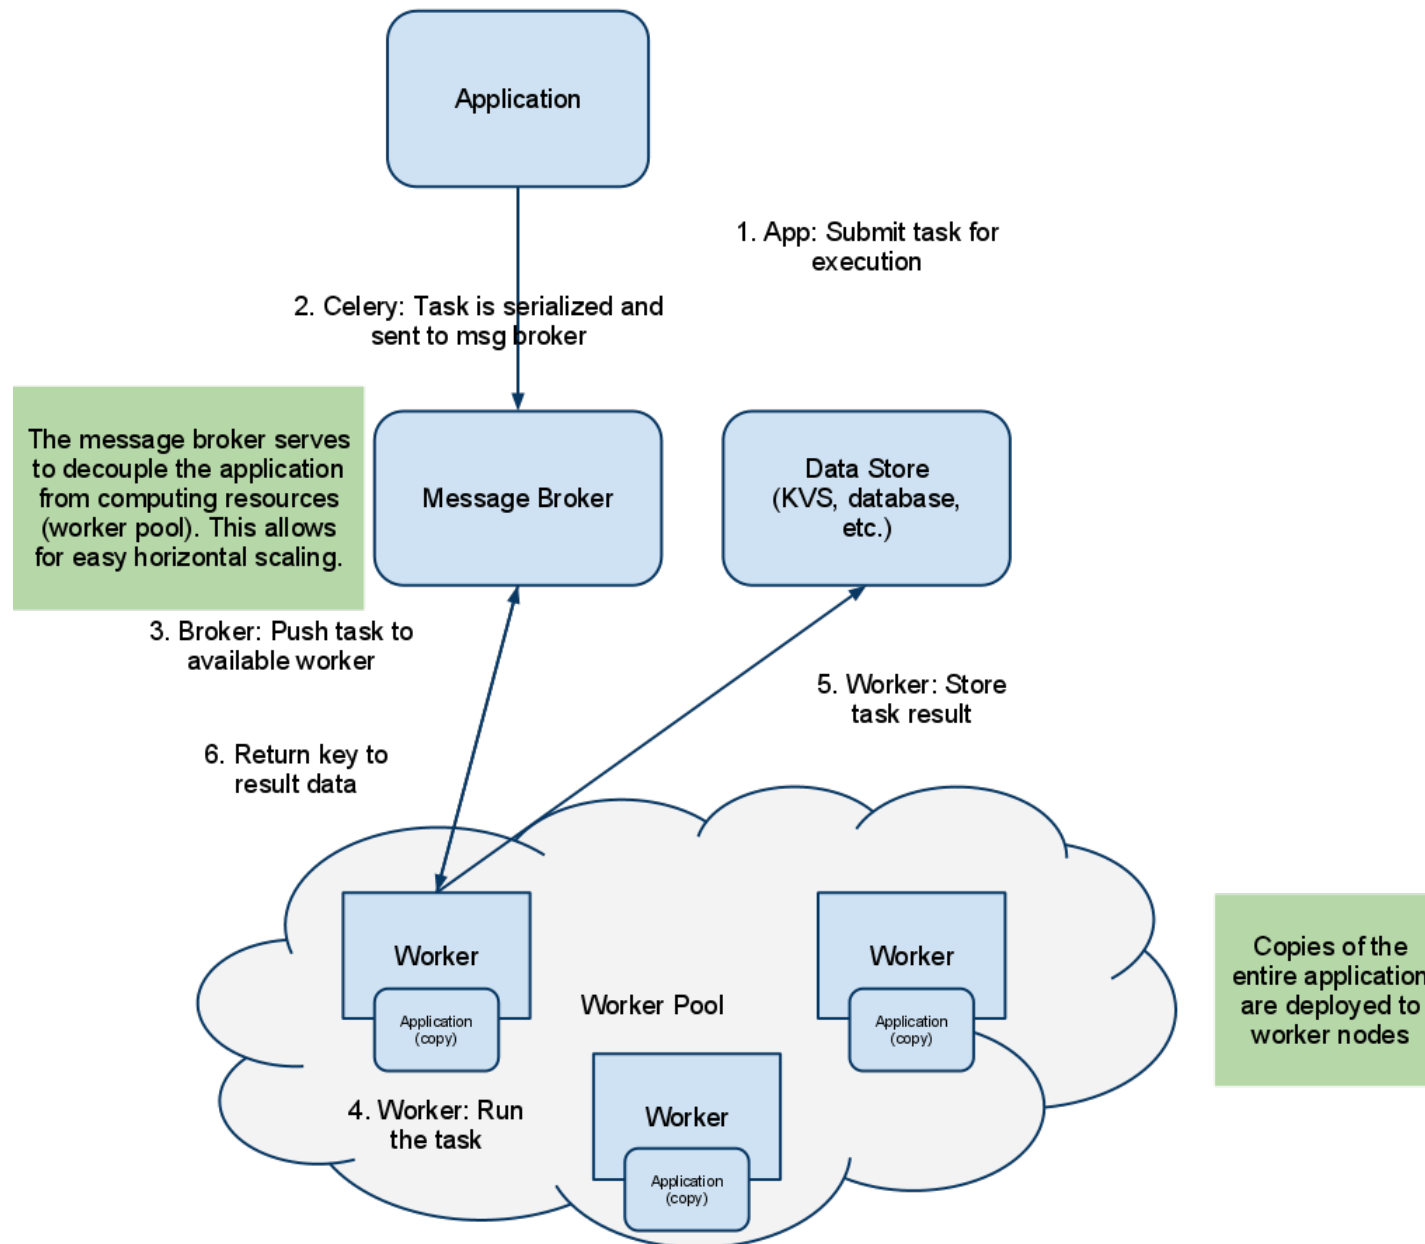
BROKER_BACKEND = "sqlalchemy"
BROKER_HOST = "sqlite:///backend.db"

# optional: "database" is the default result backend
# CELERY_RESULT_BACKEND = "database"
CELERY_RESULT_DBURI = "sqlite:///result.db"

# optional: enable verbose logging from SQLAlchemy.
# CELERY_RESULT_ENGINE_OPTIONS = {"echo": True}

CELERY_IMPORTS = ("tasks",)
```





Code Examples

<https://github.com/larsbutler/celery-examples>

(These slides are in the repo as well.)